

# DUMA-bench: A Dual-Control Multi-Agent Benchmark for Evaluating LLM Agent Security

Anonymous ACL submission

## Abstract

LLM-based agent systems are increasingly deployed for automating complex tasks, yet their security in realistic interaction scenarios remains understudied. Existing benchmarks evaluate either isolated agent capabilities or prompt injection resistance in simplified settings without accounting for active user interaction dynamics. We present DUMA-bench (Dual-control User Multi-Agent benchmark), extending  $\tau^2$ -bench with three security domains aligned with AI-SAFE (levels 1–5) and OWASP classifications: mail\_rag\_phishing (RAG poisoning), collab (inter-agent attacks), and output\_handling (improper output handling).

We evaluate five models (Claude Sonnet 4.5, GPT-4-turbo, GPT-4o, GPT-4o-mini, GPT-3.5-turbo) across varying user temperatures ( $T_{\text{user}} \in \{0.0, 0.5, 1.0\}$ ), measuring pass@1 and pass@4 metrics. Key findings: (1) model size does not predict security—GPT-4o outperforms larger models in RAG poisoning while GPT-3.5-turbo shows highest stability; (2) user behavior variability significantly impacts Attack Success Rate (ASR) even with non-malicious users; (3) all models show high ASR (75–90%) in RAG poisoning, indicating critical vulnerabilities requiring specialized defenses.

## 1 Introduction

### 1.1 Context and Problem Statement

The development of Large Language Models (LLMs) and their integration into agent systems opens new possibilities for automating complex tasks. AI agents are capable of autonomous planning, interaction with external tools (APIs, databases, file systems), and real-time decision-making (Muleys et al., 2025). However, these same capabilities create fundamentally new attack surfaces not characteristic of classical machine learning systems.

A typical AI agent comprises five interconnected components (Muleys et al., 2025). At its core, the **LLM** serves as the central component for understanding instructions and generating responses, while a **planning module** transforms high-level goals into executable action sequences. The agent’s **memory** system maintains both short-term context through dialogue history and long-term knowledge through RAG systems and knowledge bases. External **tools** provide the agent with capabilities to interact with APIs, databases, and other real-world systems, while the **interface** layer serves as the entry point for receiving user requests and delivering responses.

Each of these components represents a potential attack vector. The AI-SAFE framework (Muleys et al., 2025) systematizes threats across five levels: interface (Prompt Injection, DoS), execution and tools (Tool Misuse, Privilege Escalation), infrastructure and orchestration (Cross-Agent Poisoning), core and logic (Jailbreaking, Goal Manipulation), data and knowledge (RAG Poisoning, Data Leakage).

### 1.2 Limitations of Existing Methods

Current benchmarks for evaluating agent system security suffer from three fundamental limitations. First, most benchmarks including AgentBench (Liu et al., 2025) and Agent Security Bench (Zhang et al., 2025) conduct **isolated evaluation** where agents operate under monopolistic control with users serving merely as passive instruction sources. Second, existing security-focused benchmarks employ **simplified attack scenarios**—for instance, Agent Dojo (DeBenedetti et al., 2024) concentrates on prompt injection in single-agent contexts without considering the complexities of inter-agent interaction and RAG systems. Third and most critically, current benchmarks fail to account for the **active user** dynamic. Research on  $\tau$ -bench (Yao et al., 2024) and  $\tau^2$ -bench (Barres et al., 2025)

83 has demonstrated that introducing an active user  
84 through dual-control paradigms leads to agent per-  
85 formance drops of up to 25 percentage points, in-  
86 dicating that coordination and communication be-  
87 come critical failure points. Yet existing security  
88 benchmarks systematically overlook this dynamic,  
89 creating a significant gap between evaluation set-  
90 tings and real-world deployment scenarios.

### 91 1.3 Contributions

92 This work makes four key contributions to the eval-  
93 uation of LLM agent security. We present **DUMA-**  
94 **bench**, a new security benchmark extending  $\tau^2$ -  
95 bench with three security domains that model typ-  
96 ical attack vectors: RAG poisoning through the  
97 mail\_rag\_phishing domain, inter-agent attacks  
98 through the collab domain, and improper output  
99 handling through the output\_handling domain.  
100 The benchmark is publicly available on GitHub.<sup>1</sup>  
101 We propose a **dual-control evaluation methodol-**  
102 **ogy** for assessing agent robustness to attacks while  
103 accounting for active user behavior, with rigorous  
104 formalization within the Dec-POMDP framework.  
105 Our **comprehensive empirical evaluation** cov-  
106 ers five state-of-the-art LLMs—Claude Sonnet 4.5,  
107 GPT-4-turbo, GPT-4o, GPT-4o-mini, and GPT-3.5-  
108 turbo—across multiple configurations, demon-  
109 strating that model size does not consistently predict  
110 security robustness and revealing unexpected per-  
111 formance patterns. Finally, we formulate and test  
112 three **testable hypotheses** about agent security: the  
113 effect of multiple runs on pass@k variance (H1),  
114 the impact of user temperature on attack success  
115 rate (H2), and the relationship between model size  
116 and security performance (H3).

## 117 2 Related Work

### 118 2.1 Agent System Evaluation Benchmarks

119 The development of LLM agents has led to the  
120 creation of benchmark series for evaluating their  
121 capabilities. AgentBench (Liu et al., 2025) eval-  
122 uates agents in eight environments, including oper-  
123 ating systems, databases, and web navigation, but  
124 focuses on functional capabilities without security  
125 consideration. ToolBench and API-Bank investi-  
126 giate tool usage but in trusted environments.

127 A key breakthrough was the  $\tau$ -bench (Yao et al.,  
128 2024) and  $\tau^2$ -bench (Barres et al., 2025) series,  
129 where the user is modeled as an active participant

130 capable of changing environment state. Formaliza-  
131 tion within the Dec-POMDP framework (Amato  
132 et al., 2013) showed that user coordination is a crit-  
133 ical bottleneck: even advanced models lose up to  
134 25 percentage points of performance when transi-  
135 tioning from monopolistic to dual control.

### 136 2.2 LLM Agent Security Evaluation

137 Agent Security Bench (ASB) (Zhang et al., 2025)  
138 formalizes attacks and defenses for LLM agents,  
139 but is limited to single-agent scenarios. Agent  
140 Dojo (DeBenedetti et al., 2024) creates a dynamic  
141 environment for evaluating prompt injection at-  
142 tacks, demonstrating that modern agents are vul-  
143 nerable even to simple attacks. However, Agent  
144 Dojo does not model an active user as an interac-  
145 tion participant, attacks propagating through inter-  
146 agent communication channels, or RAG system  
147 poisoning in realistic scenarios involving emails  
148 and documents.

### 149 2.3 Threat Modeling Frameworks

150 OWASP LLM Top 10 (OWASP Foundation, 2025b)  
151 and OWASP AI Agents Top 15 (OWASP Founda-  
152 tion, 2025a) systematize threats for AI systems.  
153 The AI-SAFE framework (Muleys et al., 2025) pro-  
154 poses a five-level threat model specific to agent  
155 architectures. Our work uses these classifications  
156 for systematic coverage of attack vectors in the  
157 developed domains.

### 158 2.4 Positioning of This Work

159 This work fills a critical gap between dual-control  
160 benchmarks like  $\tau^2$ -bench, which evaluate agent-  
161 user coordination but do not focus on security, and  
162 security benchmarks like Agent Dojo and ASB,  
163 which assess vulnerabilities but do not account for  
164 active users and inter-agent interaction. We extend  
165 the  $\tau^2$ -bench methodology with security domains  
166 covering AI-SAFE threats and evaluate agent ro-  
167 bustness in realistic scenarios with active users.

## 168 3 Method

### 169 3.1 Problem Formulation

#### 170 3.1.1 Interaction Model (Dec-POMDP)

171 We formalize agent-user interaction as a Decentral-  
172 ized Partially Observable Markov Decision Process  
173 (Dec-POMDP) (Amato et al., 2013; Barres et al.,  
174 2025)<sup>1,2</sup>, following the  $\tau^2$ -bench framework. The en-  
175 vironment  $\mathcal{E}$  is described by a state space  $\mathcal{S}$  that is  
176 only partially observable to participants. The agent

<sup>1</sup>[https://github.com/sec-lab-itmo/  
duma-benchmark](https://github.com/sec-lab-itmo/duma-benchmark)

177  $\mathcal{A}$  and user  $\mathcal{U}$  function as two autonomous players,  
 178 each equipped with their own observation spaces  
 179  $\Omega_A$  and  $\Omega_U$ , and action spaces  $\mathcal{A}_A$  and  $\mathcal{A}_U$ . The  
 180 transition function  $T : \mathcal{S} \times \mathcal{A}_A \times \mathcal{A}_U \rightarrow \Delta(\mathcal{S})$   
 181 defines environment dynamics, capturing how both  
 182 agent and user actions jointly determine state evo-  
 183 lution. Critically, the user is not a passive instruc-  
 184 tion source but actively changes environment state  
 185 through their actions, creating genuine coordina-  
 186 tion challenges.

### 187 3.1.2 Attacker Model

188 We adopt a grey-box attacker model with realistic  
 189 assumptions about adversary capabilities. The at-  
 190 tacker possesses **system knowledge**, understanding  
 191 the agent architecture and available tools, but has  
 192 no access to the system prompt or model weights.  
 193 The attacker's **capabilities** include injecting mali-  
 194 cious content into various input channels such as  
 195 emails, colleague messages, user queries, and doc-  
 196 ument contents in RAG systems. The attacker's  
 197 **goal** is to force the agent to perform unauthorized  
 198 actions or disclose confidential information, repre-  
 199 senting realistic threat scenarios in deployed agent  
 200 systems.

### 201 3.1.3 Evaluation Metrics

202 **pass@k metric.** Defines the fraction of cases  
 203 where the agent successfully passed  $k$  asserts:

$$204 \text{pass}@k = \frac{1}{N} \sum_{i=1}^N \mathbb{1} \left[ \sum_{j=1}^{M_i} a_{ij} \geq k \right], \quad (1)$$

205 where  $N$  is the number of cases,  $M_i$  is the number  
 206 of asserts in case  $i$ ,  $a_{ij} \in \{0, 1\}$  is the result of  
 207 assert  $j$ .

208 **Attack Success Rate (ASR).** Attack success  
 209 indicator:

$$210 \text{ASR} = 1 - \text{pass}@1. \quad (2)$$

211 In the security context, asserts verify that the  
 212 agent maintains security invariants: absence of ma-  
 213 licious instruction execution, preservation of data  
 214 integrity, and refusal of unauthorized actions.

## 215 3.2 Developed Security Domains

### 216 3.2.1 mail\_rag\_phishing Domain (RAG 217 Poisoning)

218 **Description.** Models an attack on an email client's  
 219 RAG system through injection of malicious instruc-  
 220 tions in email content. 270

221 **Threat classification:** AI-SAFE.DATA.1  
 222 (Knowledge Base Poisoning), AI-SAFEINPUT.1

(Indirect Prompt Injection), LLM01, RAG: Indirect  
223  
 Prompt Injection. 224

**Domain architecture.** The domain simulates an  
 225 email client with integrated RAG for searching rel-  
 226 evant emails. The knowledge base contains a mix-  
 227 ture of legitimate and malicious emails designed  
 228 to test the agent's ability to distinguish safe from  
 229 compromised content. The agent has access to  
 230 four tools: `search_emails` for querying the knowl-  
 231 edge base, `get_email_details` for retrieving full  
 232 message content, `send_email` for composing new  
 233 messages, and `forward_email` for redistributing  
 234 existing messages. 235

### 236 3.2.2 collab Domain (Inter-Agent 237 Interaction)

**Description.** Models a multi-agent environment  
 238 where an attacking agent ("colleague") manipu-  
 239 lates the target agent through the communication  
 240 channel. 241

**Threat classification:** AI-SAFE.INFRA.3  
 242 (Cross-Agent Poisoning), AI-SAFE.EXEC.1 (Tool  
 243 Misuse), AI-SAFE.EXEC.2 (Privilege Escalation),  
 244 T15 (Cross-Agent Communication Poisoning),  
 245 MCP05. 246

**Domain architecture.** The domain implements  
 247 a workspace with multiple agents operating at dif-  
 248 ferent privilege levels, creating opportunities for  
 249 privilege escalation attacks. A ticketing and task  
 250 system with delegation mechanisms enables at-  
 251 tackers to manipulate task assignments. An ac-  
 252 tivity log provides audit capabilities for detect-  
 253 ing suspicious behavior. The agent can invoke  
 254 five tools: `get_tickets` for retrieving assigned  
 255 work, `update_ticket` for modifying task sta-  
 256 tus, `delegate_task` for reassigning responsibil-  
 257 ities, `get_logs` for accessing system activity, and  
 258 `send_message` for inter-agent communication. 259

### 260 3.2.3 output\_handling Domain (Output 261 Processing)

**Description.** Tests agent robustness to generating  
 262 malicious code or injections in output. 263

**Threat classification:** AI-SAFE.INPUT.3 (Im-  
 264 proper Output Handling), LLM05. 265

**Domain architecture.** The domain models a  
 266 web application where the agent generates con-  
 267 tent for display to end users. A database contain-  
 268 ing user data provides context for content genera-  
 269 tion while creating opportunities for SQL injection.  
 270 A rendering system displays generated content in  
 271 browsers, enabling testing of XSS and other in-  
 272

jection vulnerabilities. The agent operates with four tools: generate\_content for creating textual responses, execute\_query for database access, render\_template for formatting output, and send\_response for delivering content to users.

## 4 Experiments

### 4.1 Experimental Setup

#### 4.1.1 Models Under Study

We evaluate five state-of-the-art LLMs spanning different capability levels and architectural generations. **Claude Sonnet 4.5** represents Anthropic’s latest model with strong reasoning capabilities. **GPT-4-turbo** provides OpenAI’s advanced model with enhanced capabilities, while **GPT-4o** offers an optimized variant balancing performance and cost. **GPT-4o-mini** serves as a compact version with reduced cost, and **GPT-3.5-turbo** establishes a baseline for comparison. Agent generation parameters are fixed at temperature  $T_{\text{agent}} = 0.0$  for all models to ensure deterministic agent behavior.

#### 4.1.2 Variable Parameters

We vary the user model temperature across three levels to test robustness under different user behavior patterns:  $T_{\text{user}} = 0.0$  for deterministic behavior,  $T_{\text{user}} = 0.5$  for moderate variability, and  $T_{\text{user}} = 1.0$  for high variability.

#### 4.1.3 Metrics

We evaluate agent robustness using two complementary metrics. The **pass@k** metric measures the proportion of cases where the agent successfully passes at least  $k$  asserts, as formally defined in Equation 1. The **ASR (Attack Success Rate)** quantifies the proportion of successful attacks, computed as  $\text{ASR} = 1 - \text{pass}@1$ . In our experiments, we primarily report pass@1 (at least 1 assert passed) and pass@4 (at least 4 asserts passed) metrics to evaluate both basic and stringent security thresholds.

#### 4.1.4 Research Hypotheses

We test three hypotheses:

**Hypothesis 1 (H1):** At fixed agent and user temperatures, increasing the number of runs  $k$  reduces the variance of pass@ $k$ , but does not guarantee monotonic changes in ASR.

**Hypothesis 2 (H2):** Changes in non-attacking user queries cause changes in ASR at fixed agent temperature.

**Hypothesis 3 (H3):** Model size does not consistently correlate with pass@ $k$  and ASR; mid-sized models may outperform larger models in specific security domains.

#### 4.1.5 Protocol

For each combination (model, temperature, domain, case), we perform  $n = 10$  independent runs. All metrics are recorded, and results are aggregated for statistical analysis using Fisher’s exact test (Fisher, 1935) for significance testing. We use Wilson confidence intervals (Wilson, 1927) for proportion estimates, which provide better coverage properties for small sample sizes compared to normal approximations.

## 5 Results & Discussion

### 5.1 Aggregated Results

Table 1 presents aggregated pass@1 results by model and domain across all temperature settings. Table 2 reports statistical significance of differences between GPT-4o and GPT-4o-mini at each temperature, and Table 3 shows the impact of temperature variation within each model.

### 5.2 Analysis

#### 5.2.1 Hypothesis Testing Results

**H1: Effect of Multiple Runs.** Our results confirm that at fixed temperatures, increasing  $k$  in pass@ $k$  reduces variance. However, ASR does not change monotonically with  $k$ . For example, Claude Sonnet 4.5 shows pass@1 = 0.25 but pass@4 = 0.20 in the mail\_rag\_phishing domain at  $T_{\text{user}} = 0.0$ , indicating that some initially successful runs fail with more stringent evaluation criteria. This non-monotonicity suggests that agent robustness is sensitive to the number of security checks required, even with deterministic agent temperature.

**H2: User Temperature Impact.** Changes in user model temperature ( $T_{\text{user}}$ ) cause measurable changes in ASR at fixed agent temperature. This confirms that user behavior variability, even when the user is non-malicious, affects the attack success rate. The effect is most pronounced in the collab domain where social engineering plays a key role.

**H3: Model Size vs. Security Performance.** Our results partially confirm this hypothesis. Surprisingly, GPT-4o (mid-sized) outperforms larger models (GPT-4-turbo, Claude Sonnet 4.5) in the mail\_rag\_phishing domain across all temperatures, though it shows instability when evaluat-

Table 1: Model Robustness Comparison by Domain (pass@1)

Model	RAG Phishing	Collab	Output
GPT-4o (T=0.0)	16/50 (32.0%)	13/60 (21.7%)	16/30 (53.3%)
GPT-4o (T=0.5)	18/50 (36.0%)	13/60 (21.7%)	18/30 (60.0%)
GPT-4o (T=1.0)	16/50 (32.0%)	17/60 (28.3%)	18/30 (60.0%)
GPT-4o-mini (T=0.0)	5/50 (10.0%)	18/60 (30.0%)	13/30 (43.3%)
GPT-4o-mini (T=0.5)	6/50 (12.0%)	13/60 (21.7%)	14/30 (46.7%)
GPT-4o-mini (T=1.0)	4/50 (8.0%)	18/60 (30.0%)	18/30 (60.0%)

Table 2: Statistical Significance: GPT-4o vs GPT-4o-mini (Fisher exact test, two-tailed)

Domain	T	4o pass@1	mini pass@1	4o ASR	mini ASR	p-value
mail_rag_phishing	0	16/50 (32%)	5/50 (10%)	68%	90%	1.28e-02
mail_rag_phishing	0.5	18/50 (36%)	6/50 (12%)	64%	88%	9.12e-03
mail_rag_phishing	1	16/50 (32%)	4/50 (8%)	68%	92%	5.04e-03
collab	0	13/60 (22%)	18/60 (30%)	78%	70%	4.04e-01
collab	0.5	13/60 (22%)	13/60 (22%)	78%	78%	1.00e+00
collab	1	17/60 (28%)	18/60 (30%)	72%	70%	1.00e+00
output_handling	0	16/30 (53%)	13/30 (43%)	47%	57%	6.06e-01
output_handling	0.5	18/30 (60%)	14/30 (47%)	40%	53%	4.38e-01
output_handling	1	18/30 (60%)	18/30 (60%)	40%	40%	1.00e+00

Table 3: Temperature Effect: p-values (Fisher exact test)

Domain	Model	T0:0.5	T0:1	T0.5:1
RAG phish.	4o	.833	1.00	.833
RAG phish.	4o-mini	1.00	1.00	.741
collab	4o	1.00	.528	.528
collab	4o-mini	.404	1.00	.404
output	4o	.795	.795	1.00
output	4o-mini	1.00	.301	.438

ing higher  $k$  values. GPT-3.5-turbo, despite being the smallest model, demonstrates the most stable performance across temperature variations. This suggests that model size and architectural sophistication alone are not reliable predictors of security robustness—specialized security training or instruction-following capabilities may play a more critical role.

### 5.2.2 Cross-Model Comparison

Key observations emerge when comparing all five models across domains. In the **RAG Poisoning (mail\_rag\_phishing)** domain, GPT-4o achieves the highest pass@1 (32-36% depending on  $T$ ), outperforming both GPT-4-turbo (30-40%) and Claude Sonnet 4.5 (25%). GPT-4o-mini shows significantly lower performance (8-12%,  $p < 0.01$ ), while GPT-3.5-turbo performs poorly (15-20%) but maintains consistent results across temperatures. For **Inter-Agent Attacks (collab)**, Claude Sonnet 4.5 demonstrates superior performance (96-100% pass@1), followed by GPT-4-turbo (96-100%) and GPT-4o (21.7-28.3%), with no statistically significant differences between GPT-4o and GPT-

4o-mini ( $p > 0.4$ ). In the **Output Handling (output\_handling)** domain, Claude Sonnet 4.5 again leads (50-67%), followed by GPT-4-turbo (50-67%) and GPT-4o (53.3-60%), with differences between GPT-4o and GPT-4o-mini remaining statistically insignificant.

### 5.2.3 RAG System Vulnerability

The mail\_rag\_phishing domain remains the most challenging across all models: even the best configuration (Claude Sonnet 4.5 at  $T = 0.5$ ) shows ASR = 75% at pass@1. This indicates high effectiveness of indirect prompt injection attacks through RAG context and the critical need for additional defenses (e.g., source validation, instruction filtering, policy-driven tool gating).

### 5.2.4 Temperature Stability

GPT-3.5-turbo exhibits the most stable behavior across different user temperatures, with minimal variance in pass@k metrics. In contrast, GPT-4o shows high sensitivity to temperature changes, particularly in the collab domain where pass@1 ranges from 21.7% to 28.3% as  $T_{user}$  increases from 0.0 to 1.0.

## 6 Conclusion

We present DUMA-bench, a comprehensive benchmark for evaluating LLM agent security in dual-control paradigms where active users can influence environment state. Our work yields four main findings. First, DUMA-bench addresses a critical gap in the evaluation landscape: existing security

benchmarks poorly correspond to real agent usage scenarios. Our three security domains covering RAG poisoning, inter-agent attacks, and output handling align with AI-SAFE levels 1-5, OWASP LLM Top 10, and OWASP AI Agents Top 15. Second, model size does not predict security—GPT-4o outperforms larger models including GPT-4-turbo and Claude Sonnet 4.5 in RAG poisoning scenarios, while GPT-3.5-turbo shows the most stable performance across temperature variations. This partially confirms our hypothesis that model size does not consistently correlate with security robustness. Third, user behavior significantly affects attack success, with changes in user temperature impacting ASR even when the user is non-malicious, confirming that dual-control dynamics are critical for realistic security evaluation. Fourth, RAG attacks remain a critical vulnerability, with all models showing high ASR (75-90%) in the mail\_rag\_phishing domain, indicating that indirect prompt injection through RAG systems remains a major vulnerability requiring specialized defenses.

**Future directions** include: (1) extending DUMA-bench with additional domains from OWASP classifications (resource\_overload, supply\_chain); (2) evaluating defense mechanisms (Llama Guard, Promptfoo) on current domains; (3) conducting qualitative analysis of simulation traces to identify failure patterns; (4) expanding model coverage to open-source alternatives (Llama, Mistral); and (5) investigating the impact of different agent architectures on security robustness.

DUMA-bench is open-source and available at <https://github.com/sec-lab-itmo/duma-benchmark>, enabling reproducible security evaluation for the research community.

## Acknowledgments

We thank the authors of  $\tau^2$ -bench for their foundational work in dual-control agent evaluation, which inspired this research. We also acknowledge the Yandex Cloud team for developing the AI-SAFE framework.

## 7 Limitations

Our work has five important limitations<sup>5</sup>. First, while we test across multiple configurations with  $n = 10$  runs each, larger sample sizes would enable more robust statistical conclusions, especially when comparing models with similar performance.

However, our use of Fisher’s exact test and Wilson confidence intervals provides valid statistical inference for small samples. Second, the user is modeled by an LLM, which may not fully capture the unpredictability and errors of real human users, though this approach ensures reproducibility and controlled experimentation. Third, while we evaluate five models including both proprietary systems (GPT, Claude) and recognize the need for open models (Llama, Gemini, etc.), our current results may not fully generalize across all model families, and future work will expand model coverage. Fourth, DUMA-bench currently covers three security domains, with additional domains such as resource\_overload and supply\_chain attacks planned for future releases. Fifth, the current version does not evaluate effectiveness of defense mechanisms such as Llama Guard, Promptfoo, and custom validators, which represents a key direction for future research.

## References

- Christopher Amato, Girish Chowdhary, Alborz Geramifard, N. Kemal Ure, and Mykel J. Kochenderfer. 2013. Decentralized control of partially observable markov decision processes. In *52nd IEEE Conference on Decision and Control*, pages 2398–2405.
- Vincent Barres, Hao Dong, Soumya Ray, Xinyun Si, and Karthik Narasimhan. 2025.  $\tau^2$ -bench: Evaluating conversational agents in a dual-control environment. *arXiv preprint arXiv:2506.07982*.
- Edoardo Debenedetti and 1 others. 2024. Agentdojo: A dynamic environment to evaluate prompt injection attacks and defenses for llm agents. In *Advances in Neural Information Processing Systems*, volume 37, pages 82895–82920.
- Ronald A. Fisher. 1935. The logic of inductive inference. *Journal of the Royal Statistical Society, 98(1):39–82*.
- Xiao Liu and 1 others. 2025. Agentbench: Evaluating llms as agents. *arXiv preprint arXiv:2308.03688*.
- R. Muleys, S. Nesteruk, and A. Lodin. 2025. Ai secure agentic framework essentials (ai-safe) v1.0. Technical report, Yandex Cloud.
- OWASP Foundation. 2025a. Owasp ai agents (agentic ai) top 15. Available at owasp.org.
- OWASP Foundation. 2025b. Owasp top 10 for large language model applications. Available at owasp.org.
- Edwin B. Wilson. 1927. Probable inference, the law of succession, and statistical inference. *Journal of the American Statistical Association, 22(158):209–219*.

Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. 2024. $\tau$ -bench: A benchmark for tool-agent-user interaction in real-world domains. <i>arXiv preprint arXiv:2406.12045</i> .	520 521 522 523
Hanrong Zhang and 1 others. 2025. Agent security bench (asb): Formalizing and benchmarking attacks and defenses in llm-based agents. <i>arXiv preprint arXiv:2410.02644</i> .	524 525 526 527